

Hard vs Standard Pairs Report

Compares model accuracy on non-trivial ('hard') vs typical ('standard') positive and negative pairs.

Hard set = edge-case pairs (hand-curated) + uncertain pairs (decimal/integer GT sign flip).

Standard set = all remaining GT=+1 or GT=-1 pairs not flagged as hard.

Hard Positive (n=34)

Positive pairs drawn from the edge-case mask (n=30) and uncertain pairs where the integer GT = +1 (n=6).

These are synergies that are either non-obvious, conditional, or where rater scores flipped direction during consensus.

Pair List

Card A	Card B	GT	Source
Armaments	Blood for Blood	+1	edge case
Armaments	Reaper	+1	edge case
Bash	Havoc	+1	edge case
Battle Trance	Second Wind	+1	edge case
Berserk	Body Slam	+1	edge case
Berserk	True Grit	+1	edge case
Brutality	Second Wind	+1	edge case
Burning Pact	Second Wind	+1	edge case
Burning Pact	Sentinel	+1	edge case
Corruption	Battle Trance	+1	uncertain
Corruption	Power Through	+1	edge case
Corruption	Second Wind	+1	edge case
Dark Embrace	Warcry	+1	edge case
Dual Wield	Second Wind	+1	edge case
Evolve	Anger	+1	uncertain
Evolve	Havoc	+1	edge case
Evolve	Second Wind	+1	edge case
Fire Breathing	Dark Embrace	+1	edge case
Fire Breathing	Evolve	+1	edge case
Headbutt	Fire Breathing	+1	edge case
Offering	Corruption	+1	edge case
Pommel Strike	Second Wind	+1	edge case
Power Through	Brutality	+1	edge case
Rampage	Fiend Fire	+1	uncertain
Rampage	Havoc	+1	edge case, uncertain
Rampage	True Grit	+1	uncertain
Seeing Red	True Grit	+1	edge case

Sentinel	Corruption	+1	edge case, uncertain
Shockwave	Blood for Blood	+1	edge case
Shockwave	Reaper	+1	edge case
Shrug it Off	Second Wind	+1	edge case
True Grit	Sentinel	+1	edge case
True Grit	Sever Soul	+1	edge case
Warcry	Fire Breathing	+1	edge case

Model Performance

Model	n	Accuracy	Prec	Recall	F1 (target)	Macro F1	Pred -1	Pred 0	Pred +1
gpt4o	34	52.9%	100.0%	52.9%	69.2%	23.1%	2 (6%)	14 (41%)	18 (53%)
gpt54	34	50.0%	100.0%	50.0%	66.7%	22.2%	5 (15%)	12 (35%)	17 (50%)
gemini10pro	34	61.8%	100.0%	61.8%	76.4%	19.1%	3 (9%)	8 (24%)	21 (62%)
gemini15flash	34	35.3%	100.0%	35.3%	52.2%	17.4%	1 (3%)	21 (62%)	12 (35%)
gpt4omini	34	47.1%	100.0%	47.1%	64.0%	21.3%	3 (9%)	15 (44%)	16 (47%)
gpt4ominift	34	35.3%	100.0%	35.3%	52.2%	17.4%	3 (9%)	19 (56%)	12 (35%)
gemini3flash	34	70.6%	100.0%	70.6%	82.8%	27.6%	3 (9%)	7 (21%)	24 (71%)
llama4maverick	34	70.6%	100.0%	70.6%	82.8%	27.6%	1 (3%)	9 (26%)	24 (71%)

Standard Positive (n=913)

All GT=+1 pairs **not** in the hard positive set.
These are synergies that are relatively clear-cut.

Model Performance

Model	n	Accuracy	Prec	Recall	F1 (target)	Macro F1	Pred -1	Pred 0	Pred +1
gpt4o	913	71.7%	100.0%	71.7%	83.5%	27.8%	23 (3%)	235 (26%)	655 (72%)
gpt54	913	67.6%	100.0%	67.6%	80.7%	26.9%	31 (3%)	265 (29%)	617 (68%)
gemini10pro	913	80.8%	100.0%	80.8%	89.4%	14.9%	36 (4%)	135 (15%)	738 (81%)
gemini15flash	913	50.1%	100.0%	50.1%	66.7%	22.2%	23 (3%)	433 (47%)	457 (50%)
gpt4omini	913	70.1%	100.0%	70.1%	82.4%	27.5%	15 (2%)	258 (28%)	640 (70%)
gpt4ominift	913	52.8%	100.0%	52.8%	69.1%	23.0%	17 (2%)	414 (45%)	482 (53%)
gemini3flash	913	76.2%	100.0%	76.2%	86.5%	28.8%	14 (2%)	203 (22%)	696 (76%)
llama4maverick	913	76.8%	100.0%	76.8%	86.9%	29.0%	18 (2%)	194 (21%)	701 (77%)

Hard Negative (n=9)

Negative pairs drawn from the edge-case mask (n=4) and uncertain pairs where the integer GT = -1 (n=6).

Small sample (n=9). Treat aggregate metrics as illustrative.

Pair List

Card A	Card B	GT	Source
Battle Trance	Warcry	-1	edge case
Burning Pact	Fiend Fire	-1	uncertain
Dual Wield	Fiend Fire	-1	uncertain
Immolate	True Grit	-1	uncertain
Offering	Sever Soul	-1	edge case
Offering	True Grit	-1	edge case, uncertain
Power Through	True Grit	-1	uncertain
True Grit	Armaments	-1	edge case
Wild Strike	True Grit	-1	uncertain

Model Performance

Model	n	Accuracy	Prec	Recall	F1 (target)	Macro F1	Pred -1	Pred 0	Pred +1
gpt4o	9	33.3%	100.0%	33.3%	50.0%	16.7%	3 (33%)	2 (22%)	4 (44%)
gpt54	9	22.2%	100.0%	22.2%	36.4%	12.1%	2 (22%)	4 (44%)	3 (33%)
gemini10pro	9	11.1%	100.0%	11.1%	20.0%	6.7%	1 (11%)	2 (22%)	6 (67%)
gemini15flash	9	22.2%	100.0%	22.2%	36.4%	12.1%	2 (22%)	5 (56%)	2 (22%)
gpt4omini	9	22.2%	100.0%	22.2%	36.4%	18.2%	2 (22%)	0 (0%)	7 (78%)
gpt4ominift	9	11.1%	100.0%	11.1%	20.0%	6.7%	1 (11%)	6 (67%)	2 (22%)
gemini3flash	9	22.2%	100.0%	22.2%	36.4%	12.1%	2 (22%)	2 (22%)	5 (56%)
llama4maverick	9	33.3%	100.0%	33.3%	50.0%	16.7%	3 (33%)	1 (11%)	5 (56%)

Standard Negative (n=116)

All GT=-1 pairs **not** in the hard negative set.
These are anti-synergies that are relatively clear-cut.

Model Performance

Model	n	Accuracy	Prec	Recall	F1 (target)	Macro F1	Pred -1	Pred 0	Pred +1
gpt4o	116	31.0%	100.0%	31.0%	47.4%	15.8%	36 (31%)	68 (59%)	12 (10%)
gpt54	116	35.3%	100.0%	35.3%	52.2%	17.4%	41 (35%)	59 (51%)	16 (14%)
gemini10pro	116	28.4%	100.0%	28.4%	44.3%	11.1%	33 (28%)	28 (24%)	54 (47%)
gemini15flash	116	25.9%	100.0%	25.9%	41.1%	13.7%	30 (26%)	64 (55%)	22 (19%)
gpt4omini	116	22.4%	100.0%	22.4%	36.6%	12.2%	26 (22%)	55 (47%)	35 (30%)
gpt4ominift	116	16.4%	100.0%	16.4%	28.1%	9.4%	19 (16%)	81 (70%)	16 (14%)
gemini3flash	116	50.0%	100.0%	50.0%	66.7%	22.2%	58 (50%)	48 (41%)	10 (9%)
llama4maverick	116	31.9%	100.0%	31.9%	48.4%	16.1%	37 (32%)	54 (47%)	25 (22%)

Summary: Hard vs Standard

Direct comparison of accuracy on hard vs standard sub-sets for each polarity.

Positive pairs

Model	Hard acc (n=34)	Standard acc (n=913)	Drop
gpt4o	52.9%	71.7%	+18.8pp
gpt54	50.0%	67.6%	+17.6pp
gemini10pro	61.8%	80.8%	+19.1pp
gemini15flash	35.3%	50.1%	+14.8pp
gpt4omini	47.1%	70.1%	+23.0pp
gpt4ominift	35.3%	52.8%	+17.5pp
gemini3flash	70.6%	76.2%	+5.6pp
llama4maverick	70.6%	76.8%	+6.2pp

Negative pairs

Model	Hard acc (n=9)	Standard acc (n=116)	Drop
gpt4o	33.3%	31.0%	-2.3pp
gpt54	22.2%	35.3%	+13.1pp
gemini10pro	11.1%	28.4%	+17.3pp
gemini15flash	22.2%	25.9%	+3.6pp
gpt4omini	22.2%	22.4%	+0.2pp
gpt4ominift	11.1%	16.4%	+5.3pp
gemini3flash	22.2%	50.0%	+27.8pp
llama4maverick	33.3%	31.9%	-1.4pp