

Uncertain Cards Report

Pairs where the decimal GT and integer GT disagree in *direction* — i.e. the rater consensus flipped the sign during aggregation. These pairs represent genuine annotation uncertainty and are used as a diagnostic lens on model behaviour under ambiguity.

Accuracy and **F1** are computed against the integer GT (the final consensus label), since that is what models were trained to predict. The decimal GT determines *which* pairs enter each mask.

Mask: half_flipped

Decimal GT = ± 0.5 (borderline), but the integer consensus GT went to the **opposite sign** (e.g. decimal +0.5 $\rightarrow$ integer -1). These are pairs where raters leaned one way individually but the final vote reversed direction.

Total pairs: 11

Pair List

Card A	Card B	Dec GT	Int GT	Direction
Burning Pact	Fiend Fire	+0.5	-1	pos $\rightarrow$ neg
Corruption	Battle Trance	-0.5	+1	neg $\rightarrow$ pos
Dual Wield	Fiend Fire	+0.5	-1	pos $\rightarrow$ neg
Evolve	Anger	-0.5	+1	neg $\rightarrow$ pos
Immolate	True Grit	+0.5	-1	pos $\rightarrow$ neg
Offering	True Grit	+0.5	-1	pos $\rightarrow$ neg
Power Through	True Grit	+0.5	-1	pos $\rightarrow$ neg
Rampage	Havoc	-0.5	+1	neg $\rightarrow$ pos

Card A	Card B	Dec GT	Int GT	Direction
Rampage	True Grit	-0.5	+1	neg → pos
Sentinel	Corruption	-0.5	+1	neg → pos
Wild Strike	True Grit	+0.5	-1	pos → neg

Model Accuracy on These Pairs

Model	n	Accuracy	Macro F1	Pred -1	Pred 0	Pred +1
gpt4o	11	0.0%	0.0%	2 (18%)	5 (45%)	4 (36%)
gpt54	11	9.1%	7.4%	1 (9%)	6 (55%)	4 (36%)
gemini10pro	11	18.2%	12.1%	1 (9%)	4 (36%)	6 (55%)
gemini15flash	11	18.2%	17.9%	1 (9%)	7 (64%)	3 (27%)
gpt4omini	11	18.2%	15.6%	1 (9%)	4 (36%)	6 (55%)
gpt4ominift	11	0.0%	0.0%	1 (9%)	9 (82%)	1 (9%)
gemini3flash	11	9.1%	6.7%	1 (9%)	5 (45%)	5 (45%)

Notes

⚠ Small sample (n=11). Treat accuracy figures as illustrative, not statistically reliable. Individual pair results shown above are more informative than aggregate metrics here.

Mask: solid_flipped

Decimal GT = ±1.0/±1.5/±2.0 (clear/strong), but the integer consensus GT went to the **opposite sign**.

These are the most surprising mismatches: pairs that scored clearly in one direction on the decimal scale but were labelled the other way in the final integer GT.

Total pairs: 1

Pair List

Card A	Card B	Dec GT	Int GT	Direction
Rampage	Fiend Fire	-1.0	+1	neg → pos

Model Accuracy on These Pairs

Model	n	Accuracy	Macro F1	Pred -1	Pred 0	Pred +1
gpt4o	1	0.0%	0.0%	0 (0%)	1 (100%)	0 (0%)
gpt54	1	0.0%	0.0%	1 (100%)	0 (0%)	0 (0%)
gemini10pro	1	100.0%	100.0%	0 (0%)	0 (0%)	1 (100%)
gemini15flash	1	100.0%	100.0%	0 (0%)	0 (0%)	1 (100%)
gpt4omini	1	0.0%	0.0%	0 (0%)	1 (100%)	0 (0%)
gpt4ominift	1	100.0%	100.0%	0 (0%)	0 (0%)	1 (100%)
gemini3flash	1	0.0%	0.0%	1 (100%)	0 (0%)	0 (0%)

Notes

⚠ Small sample (n=1). Treat accuracy figures as illustrative, not statistically reliable. Individual pair results shown above are more informative than aggregate metrics here.

Generated by `uncertain_cards.py`